

EU AI Act Compliance Documentation — AI Career LaunchPad

Marco Martins

File: [eu_ai_act_compliance.pdf](#)

Project: AI Career LaunchPad

Builder: Marco Martins

Context: Ironhack AI Consulting & Integration Bootcamp — Final Capstone Project

TABLE OF CONTENTS

<i>EU AI Act Compliance Documentation — AI Career LaunchPad</i>	<i>1</i>
<i>1. Risk Classification</i>	<i>3</i>
<i>2. Classification Reasoning</i>	<i>3</i>
<i>Step 1 — Is the system prohibited? (Unacceptable Risk)</i>	<i>3</i>
<i>Step 2 — Is the system High Risk under Annex III?</i>	<i>3</i>
<i>Step 3 — Does the system trigger Limited Risk obligations?</i>	<i>4</i>
<i>3. Mandatory Requirements Summary (Limited Risk)</i>	<i>5</i>
3.1 Disclosure that the system is AI	5
3.2 Disclosure of AI-generated content	5
3.3 Human oversight	5
3.4 Accuracy and robustness	6
3.5 Cybersecurity	6
<i>4. AI Manifesto — Design Philosophy and Human Oversight</i>	<i>6</i>
Do no harm. Always.	7
Human enhancement over human replacement.	7
Human judgement over automated decisions.	7
Critical thinking over blind automation.	8
Transparency over convenience.	8
Privacy over data hunger.	8
Wellbeing over efficiency.	8
Positive impact over impressive technology.	9
<i>5. Conformity Assessment Summary</i>	<i>9</i>
What the system does	9

<i>Risk class and basis for classification</i>	9
<i>Applicable obligations and how the design addresses them</i>	10
<i>Gaps not yet addressed and resolution path</i>	10
6. Technical Documentation Outline	11

1. Risk Classification

Classification: Limited Risk

AI Career LaunchPad is classified as a Limited Risk AI system under the EU AI Act (Regulation (EU) 2024/1689).

*The primary obligation triggered at this risk level is **transparency**: users must be informed they are interacting with an AI system, and the nature and limitations of AI-generated outputs must be clearly communicated at the point of use.*

2. Classification Reasoning

The EU AI Act establishes four risk tiers: Unacceptable Risk (prohibited), High Risk (Annex III), Limited Risk (transparency obligations), and Minimal Risk (no specific obligations). The classification of AI Career LaunchPad proceeds as follows.

Step 1 — Is the system prohibited? (Unacceptable Risk)

Prohibited systems include social scoring by public authorities, real-time biometric surveillance, and systems that exploit psychological vulnerabilities to manipulate behaviour.

AI Career LaunchPad does none of these. It is a privately operated career advisory tool. It does not score individuals for access to public services, does not use biometric data, and does not manipulate behaviour. Unacceptable Risk does not apply.

Step 2 — Is the system High Risk under Annex III?

Annex III lists specific high-risk categories. The categories most relevant to examine for a career advisory tool are:

***Annex III, Point 4 — Employment, workers management and access to self-employment:** This covers AI systems used for recruitment, CV screening, promotion decisions, task allocation, and monitoring of workers' performance. It applies to systems that make or materially influence employment decisions about identifiable individuals within an employment relationship or hiring process.*

AI Career LaunchPad does not make employment decisions. It does not screen CVs on behalf of an employer. It does not rank candidates for a specific vacancy. It does not operate within a hiring pipeline. It is a self-directed career development tool used by the individual about themselves — not by an employer about a candidate. The output is advisory: a Career Compass report and an application package that the individual chooses whether and how to use.

The key distinction: Annex III Point 4 targets systems that sit on the employer side of the hiring relationship and influence decisions made about workers or candidates. LaunchPad sits entirely on the candidate side and produces outputs that inform the candidate's own decisions. This is a fundamentally different relationship and a fundamentally different risk profile.

Annex III, Point 5 — Access to essential private services and public services: This covers credit scoring, insurance risk assessment, and similar gatekeeping systems. LaunchPad does not gate access to any service. Does not apply.

Annex III, Point 6 — Law enforcement and biometric categorisation: Not applicable.

Conclusion: High Risk does not apply. The system does not fall within any Annex III category. The employment-adjacent argument — that any AI system producing personalised assessments about individuals in a career context could attract scrutiny — does not survive contact with the text of the Act. The Act is explicit that High Risk designation under Point 4 requires the system to be used in an employment or recruitment decision context by an operator in a position of authority over the subject. LaunchPad is used by the subject about themselves.

Step 3 — Does the system trigger Limited Risk obligations?

Yes. Article 50 of the EU AI Act establishes transparency obligations for AI systems that interact with natural persons or generate content that could be mistaken for human-produced output.

AI Career LaunchPad generates personalised text — positioning statements, cover letters, interview guidance — that a user may present to third parties as their own work. The transparency obligation applies: users must be informed that the outputs are AI-generated, and the system must not be designed to obscure this fact.

Classification confirmed: Limited Risk.

3. Mandatory Requirements Summary (Limited Risk)

Under Article 50, the following transparency obligations apply:

Disclosure that the system is AI

Obligation: Users must be informed they are interacting with an AI system.

How the design addresses it: Every output page carries the AI Career LaunchPad branding with explicit reference to AI generation. The Career Compass report includes a transparency notice on the profile match tab: "This profile match is generated by an AI system based on your inputs. It is a starting point for reflection, not a definitive career verdict. Use your own judgement." The Application Engine output includes: "Review all content carefully before use. Do not submit the cover letter without personalising it."

Gap at POC stage: The transparency notice is present in the output but not at the point of data collection (the form trigger). This must be added before MVP deployment.

Disclosure of AI-generated content

Obligation: Where AI generates content that a person may present as their own (cover letters, positioning statements), the AI origin must be disclosed.

How the design addresses it: The cover letter tab includes an explicit instruction that the letter is AI-generated and must be reviewed and personalised before use. The elevator pitch section of the Career Compass carries the label "AI suggestion — make it yours" and includes the note: "This is a starting point. The best elevator pitch is the one you genuinely own."

Human oversight

Obligation: Limited Risk systems are not subject to the same mandatory human oversight mechanisms as High Risk systems, but the spirit of the Act — and the GDPR's restrictions on solely automated decision-making — requires that consequential outputs are not presented as final or binding.

How the design addresses it: No output from AI Career LaunchPad is presented as a final decision. Every section of both reports uses advisory language. The scoring system produces a ranked fit, not a pass/fail gate. The system explicitly invites the user to exercise their own judgement at every stage.

Gap at POC stage: There is no formal human-in-the-loop mechanism beyond the disclaimer text. For the B2C coaching model, the coaching session itself serves as the

human oversight layer. For the B2B operator model, this must be formalised — a trained facilitator or career advisor should review Career Compass outputs with participants before they act on them.

Accuracy and robustness

***Obligation:** The system should not produce outputs that are misleading or that could cause harm through inaccuracy.*

***How the design addresses it:** The extraction prompts are explicitly designed for factual grounding — the system is instructed to extract only what is in the source documents, not to infer or invent. The profile matching node includes a programmatic validation layer (Validate Profile Matching) that corrects scoring arithmetic independently of the LLM. Gap analysis is mandatory on every profile match — the system cannot produce a report that shows only strengths.*

***Gap at POC stage:** No systematic accuracy testing has been conducted against a diverse candidate population. The system has been tested with one user profile. Accuracy validation across varied backgrounds, industries, and career lengths is required before MVP deployment.*

Cybersecurity

***Obligation:** AI systems must be reasonably protected against adversarial manipulation that could alter their behaviour or outputs.*

***How the design addresses it:** The system does not accept free-text input that is directly injected into LLM prompts without structure. Inputs are PDF documents processed by an extraction node before reaching the LLM. The n8n instance uses credential-based API key storage. No candidate data is stored beyond the session.*

***Gap at POC stage:** No formal security review has been conducted. Prompt injection via crafted PDF content is a theoretical attack vector. This requires assessment before MVP deployment.*

4. AI Manifesto — Design Philosophy and Human Oversight

The following principles govern the design of AI Career LaunchPad. Each principle is mapped to the corresponding EU AI Act obligation it addresses.

These are not aspirational statements. They are active design constraints — principles that shaped specific decisions about what the system does, what it refuses to do, and how it presents its outputs.

Do no harm. Always.

Mapped to: Article 9 (Risk Management), Article 50 (Transparency obligations)

The system is designed so that no output can cause direct harm by being acted on uncritically. Every report includes mandatory gap analysis. The profile matching system is architecturally incapable of returning a report with no weaknesses — the gap array is enforced at prompt level and validated in code. The cover letter output explicitly instructs the user not to submit it without personalisation.

Human enhancement over human replacement.

Mapped to: Article 14 (Human Oversight), Recital 47 (AI as a tool supporting human decision-making)

AI Career LaunchPad does not make career decisions. It enhances the candidate's ability to make their own informed decisions by surfacing structure and language they may not have access to on their own. The system produces inputs to a human decision process — it does not substitute for one. In the B2C coaching model, the report is explicitly designed as a preparation tool for the coaching conversation, not a replacement for it.

Human judgement over automated decisions.

Mapped to: Article 14 (Human Oversight), GDPR Article 22 (Automated individual decision-making)

No output from this system is presented as a binding or final determination. Scores are advisory. Profile matches are starting points. The elevator pitch is labelled "AI suggestion — make it yours." The Career Compass report opens with a transparency notice that explicitly names it as a starting point for reflection, not a career verdict.

Critical thinking over blind automation.

Mapped to: Article 13 (Transparency and provision of information), Article 50 (Transparency obligations)

The system is designed to provoke reflection, not to provide answers. The SWOT analysis surfaces weaknesses and threats. The gap analysis names specific deficits and explains how to talk about them. The concern response section directly addresses the candidate's own stated anxieties. The report is not a validation exercise — it is a structured challenge.

Transparency over convenience.

Mapped to: Article 13 (Transparency), Article 50 (Transparency obligations)

It would be more convenient to remove transparency notices from the output — they interrupt the reading experience. They are present anyway. It would be more convenient to present the cover letter as ready to send. It is not presented that way. Every section where AI authorship could be mistaken for human authorship is explicitly labelled.

Privacy over data hunger.

Mapped to: GDPR Article 5 (Data minimisation), Article 25 (Data protection by design)

The system processes the minimum data required to produce the output. At POC stage, no candidate data is stored beyond the active session. The system does not collect data it does not need. It does not retain outputs. It does not build candidate profiles over time. See [gdpr_documentation.md](#) for the full data processing register.

Wellbeing over efficiency.

Mapped to: Article 9 (Risk Management), Recital 43 (Protection of vulnerable persons)

The system handles sensitive self-reported data — career anxieties, identity concerns, unspoken fears about mispricing. The prompts are designed to surface and address these concerns directly rather than suppress them for the sake of a more positive output. The unspoken concern section of the Career Compass explicitly names the candidate's stated anxiety and provides a direct, honest response — not a reassurance.

Positive impact over impressive technology.

Mapped to: Recital 5 (AI in the public interest), Article 69 (Codes of conduct)

The technology choices in this system are deliberately conservative. gpt-4o is used where output quality is critical. gpt-4o-mini is used where it is sufficient. The knowledge base is hardcoded rather than embedded in a vector database — because for this POC, the simpler approach produces better results and is easier to audit. Complexity is not a feature. The measure of success is whether a candidate lands a better role, not whether the architecture is impressive.

5. Conformity Assessment Summary

System name: AI Career LaunchPad Version assessed: POC v1.0 (WF1 Career Compass + WF3 Application Engine) Assessment date: April 2026 Assessed by: Marco Martins (builder)

What the system does

AI Career LaunchPad is a two-workflow AI pipeline that takes career documents and self-reflection inputs from an individual professional and produces two outputs: a personalised Career Compass report (profile matching, positioning strategy, 90-day next steps) and a job-specific Application Engine report (fit analysis, cover letter, interview preparation). Both outputs are advisory. The system does not make employment decisions, does not operate within a hiring pipeline, and does not share candidate data with employers or third parties.

Risk class and basis for classification

Limited Risk. The system does not fall within any Annex III High Risk category. The employment-adjacent argument does not apply because the system operates on the candidate side of the career transition, not the employer side. The primary obligations triggered are transparency obligations under Article 50.

Applicable obligations and how the design addresses them

<i>Obligation</i>	<i>Status</i>	<i>Evidence</i>
<i>Disclosure that system is AI</i>	<i>Partially met</i>	<i>Transparency notice present in report output. Not yet present at point of data collection.</i>
<i>Disclosure of AI-generated content</i>	<i>Met</i>	<i>Cover letter and elevator pitch explicitly labelled as AI-generated with instruction to personalise.</i>
<i>Advisory framing — no binding decisions</i>	<i>Met</i>	<i>All scoring and matching outputs use advisory language. No pass/fail gates.</i>
<i>Data minimisation</i>	<i>Met at POC stage</i>	<i>No data stored beyond session. No third-party data sharing beyond OpenAI API.</i>
<i>Accuracy safeguards</i>	<i>Partially met</i>	<i>Programmatic validation of scoring arithmetic. Gap analysis mandatory. No population-level accuracy testing yet.</i>
<i>Cybersecurity</i>	<i>Partially met</i>	<i>Credential-based API key storage. No formal security review conducted.</i>

Gaps not yet addressed and resolution path

Gap 1 — No transparency notice at point of data collection Resolution: Add a clear AI disclosure statement to the Form Trigger page before users upload documents. Implement before MVP launch. Target: pre-launch, Sprint 1.

Gap 2 — No formal human oversight mechanism beyond disclaimer text Resolution: For the B2C coaching model, formalise the coaching session as the mandatory human review stage — the Career Compass output is not acted on until reviewed with a qualified coach. For the B2B operator model, require a trained facilitator to conduct a group debrief of

Career Compass outputs with each cohort. Document this as a procedural control. Target: operational procedure defined before first operator licence signed.

Gap 3 — No audit logging Resolution: *Implement lightweight audit logging of form submissions (timestamp, completion status, no personal data) in MVP. Add structured logging of LLM node errors and validation failures. Target: Sprint 2.*

Gap 4 — No population-level accuracy validation Resolution: *Conduct structured user testing with a minimum of 20 diverse candidate profiles before commercial launch. Collect satisfaction scores and profile match accuracy ratings. Use results to calibrate prompts. Target: beta cohort (Ironhack colleagues) before MVP release.*

Gap 5 — No formal security review Resolution: *Commission a basic security review of the n8n instance and webhook endpoints before MVP deployment. Address prompt injection risk via PDF content. Target: pre-launch, Sprint 2.*

6. Technical Documentation Outline

The following table of contents represents the full technical documentation package that would be required for a formally compliant Limited Risk AI system. At POC stage, the documents marked as existing are complete. Documents marked as planned are required before commercial deployment.

<i>Document</i>	<i>Description</i>	<i>Status</i>
<i>System description</i>	<i>What the system does, inputs, outputs, intended use, out-of-scope uses</i>	<i>Exists — use_case_definition.md</i>
<i>Risk classification record</i>	<i>Classification reasoning, Annex III analysis, conclusion</i>	<i>Exists — this document, Section 2</i>
<i>Conformity assessment</i>	<i>Obligations, compliance status, gaps, resolution plan</i>	<i>Exists — this document, Section 5</i>

<i>Document</i>	<i>Description</i>	<i>Status</i>
<i>Data flow map</i>	<i>Personal data collected, processed, stored, transferred</i>	<i>Planned — gdpr_documentation.md</i>
<i>Processing activities register</i>	<i>Legal basis, retention, third-party recipients per data type</i>	<i>Planned — gdpr_documentation.md</i>
<i>DPIA</i>	<i>Data Protection Impact Assessment for highest-risk processing</i>	<i>Planned — gdpr_documentation.md</i>
<i>Transparency notice — form</i>	<i>AI disclosure at point of data collection</i>	<i>Planned — Sprint 1</i>
<i>Transparency notice — output</i>	<i>AI disclosure in report output</i>	<i>Exists — embedded in HTML output</i>
<i>Human oversight procedure</i>	<i>Formal description of the human review stage in B2C and B2B models</i>	<i>Planned — before first operator licence</i>
<i>Accuracy validation report</i>	<i>Results of population-level testing across diverse candidate profiles</i>	<i>Planned — beta cohort phase</i>
<i>Security review</i>	<i>Assessment of n8n instance, webhook endpoints, prompt injection risk</i>	<i>Planned — Sprint 2</i>
<i>Incident response procedure</i>	<i>What to do if a data breach, accuracy failure, or harmful output occurs</i>	<i>Planned — MVP phase</i>

<i>Document</i>	<i>Description</i>	<i>Status</i>
<i>Model cards</i>	<i>Description of OpenAI models used, known limitations, version pinning policy</i>	<i>Planned — MVP phase</i>
<i>Audit log specification</i>	<i>What is logged, retention period, access controls</i>	<i>Planned — Sprint 2</i>

AI Career LaunchPad — Ironhack AI Consulting & Integration Bootcamp — Week 9 Capstone