

# Benchmark Audit

Scenario: HR Recruitment Agency — Automated Job Description Generator  
Consultant: Marco Martins  
Date: April 2026

---

## TABLE OF CONTENTS

|                                                           |   |
|-----------------------------------------------------------|---|
| 1. Use Case Overview                                      | 2 |
| 1.1 Client & Context                                      | 2 |
| 1.2 Objective                                             | 2 |
| 1.3 Key Requirements                                      | 2 |
| 1.4 Key Failure Modes                                     | 2 |
| 1.5 Evaluation Strategy Summary                           | 3 |
| 2. Benchmark Evaluations                                  | 3 |
| 2.1 IFEval (Instruction Following Evaluation)             | 3 |
| 2.2 TruthfulQA                                            | 4 |
| 2.3 BOLD (Bias in Open-Ended Language Generation Dataset) | 5 |

# 1. Use Case Overview

## 1.1 Client & Context

Our client is a mid-size recruitment agency that operates across multiple industries. The agency receives structured briefs from client companies — typically a list of bullet points covering role title, key responsibilities, required skills, seniority level, and team context — and is responsible for turning these into polished, publish-ready job descriptions. Currently, this process is handled manually by HR consultants and takes between 30 and 90 minutes per role.

## 1.2 Objective

The agency wants to automate the first draft of job descriptions using a large language model. The goal is not to replace human review entirely, but to reduce turnaround time significantly, allowing consultants to focus on editing and quality control rather than writing from scratch.

## 1.3 Key Requirements

The LLM-generated job descriptions must meet the following criteria:

- **Faithfulness:** The output must accurately reflect the client brief. No skills, responsibilities, or requirements should be added that were not in the original input.
- **Completeness:** All information provided in the brief must be present in the output. Nothing critical should be omitted.
- **Professional tone:** The output must be well-structured, clearly written, and ready to publish without significant editing.
- **Inclusivity:** The language must be bias-free and inclusive. Gendered pronouns, exclusionary language, or role stereotypes must be avoided.
- **Adaptability:** The model must be able to handle briefs from different industries (tech, finance, healthcare, retail) and seniority levels (junior, senior, executive).

## 1.4 Key Failure Modes

The following failure modes have been identified as the most critical risks:

- **Hallucination:** The model invents skills or requirements not present in the brief (e.g., adding "5 years of Python experience" when the brief only mentions "programming skills").
- **Omission:** The model drops critical information from the brief, such as seniority level or key responsibilities.

- **Gendered or exclusionary language:** The model uses "he/him" pronouns, gender-coded adjectives (e.g., "aggressive", "nurturing"), or language that may discourage certain demographics from applying.
- **Generic output:** The model produces a copy-paste sounding job description that does not reflect the specific role or company culture described in the brief.
- **Inconsistent tone:** The output is unprofessional, too casual, or inconsistent in register across sections.

## 1.5 Evaluation Strategy Summary

Given these requirements and failure modes, the evaluation strategy focuses on three dimensions: **instruction following** (does the model faithfully follow the brief?), **bias and inclusivity** (is the language free from gender and demographic bias?), and **writing quality** (is the output professional, complete, and well-structured?). The three benchmarks audited below were selected to cover each of these dimensions.

---

## 2. Benchmark Evaluations

---

### 2.1 IFEval (Instruction Following Evaluation)

Benchmark Name: IFEval

Year: 2023

Source: Zhou et al., 2023 — <https://arxiv.org/abs/2311.07911>

Dataset: <https://huggingface.co/datasets/google/IFEval>

Why it seemed relevant:

Our use case requires the model to follow a structured client brief and produce output that respects multiple constraints simultaneously — tone, structure, length, and inclusivity. IFEval was designed to test whether a model can follow verifiable output constraints, such as formatting rules, word limits, and style requirements. This maps directly to what a JD generator must do when given a brief with specific requirements.

Contamination risk:

- High — Model definitely saw this during training
- Explanation: IFEval was published in late 2023 and is one of the core benchmarks on the Hugging Face Open LLM Leaderboard. Most frontier models (GPT-4, Claude, Gemini) were almost certainly trained on or fine-tuned against IFEval-style constraints, leading to inflated performance metrics. Its 25-constraint taxonomy is well-known and widely used in synthetic training data generation.

Saturation risk:

- High — Many models achieve near-perfect scores

- Explanation: Top models now score above 0.90 on IFEval (e.g., Qwen3.5-27B at 0.950, Kimi K2.5 at 0.940). The average across 61 evaluated models is 0.842, indicating the benchmark no longer effectively differentiates between frontier models.

**Format:**

- Free-form text (constrained generation with verifiable rules)

**Verdict:**

- Reject it
  - Explanation: High contamination and high saturation mean IFEval scores are unreliable for our use case. A model that scores 0.95 on IFEval may still hallucinate skills in a JD or use exclusionary language — the benchmark simply does not test for these failures. Its constraints are also generic (word counts, formatting) and do not cover domain-specific requirements such as faithfulness to a professional brief.
- 

## 2.2 TruthfulQA

**Benchmark Name:** TruthfulQA

**Year:** 2021

**Source:** Lin et al., 2022 — <https://arxiv.org/abs/2109.07958>

**Dataset:** [https://huggingface.co/datasets/truthful\\_qa](https://huggingface.co/datasets/truthful_qa)

**Why it seemed relevant:**

One of our key failure modes is hallucination — the model inventing skills, requirements, or qualifications not present in the client brief. TruthfulQA was designed to measure a model's tendency to produce confident but factually incorrect outputs, which seemed relevant to assessing whether a model would fabricate content when generating job descriptions.

**Contamination risk:**

- High — Model definitely saw this during training
- Explanation: TruthfulQA was published in 2021 and has been publicly available since then. Multiple studies have confirmed contamination: its 817 questions are widely circulated, and research has shown that removing contaminated examples from evaluation sets produces significant accuracy drops, suggesting scores reflect memorisation as much as capability.

**Saturation risk:**

- High — Many models achieve near-perfect scores
- Explanation: TruthfulQA has been reported as saturated by recent analysis, with top models clustering at high accuracy. The benchmark also contains some incorrect gold answers and metrics that excessively penalise models, further reducing its reliability as a differentiator.

**Format:**

- Multiple Choice (MC1, MC2, MC3 variants) and free-form generation

**Verdict:**

- Reject it
  - Explanation: Beyond contamination and saturation issues, TruthfulQA measures general knowledge falsehoods — not document-grounded faithfulness. Our hallucination concern is different: we want to know if the model invents content relative to a provided brief, not whether it repeats common misconceptions. TruthfulQA does not test this form of hallucination and is therefore not fit for purpose.
- 

## 2.3 BOLD (Bias in Open-Ended Language Generation Dataset)

**Benchmark Name:** BOLD

**Year:** 2021

**Source:** Dhamala et al., FAccT 2021 — <https://arxiv.org/abs/2101.11718>

**Dataset:** <https://huggingface.co/datasets/AlexaAI/bold> | <https://github.com/amazon-science/bold>

**Why it seemed relevant:**

A key requirement of our JD generator is that output must be inclusive and bias-free — no gendered language, no exclusionary tone. BOLD was designed precisely to measure social biases in open-ended text generation across five domains: profession, gender, race, religion, and political ideology. The profession and gender domains are directly relevant to job description generation, where models may default to gendered pronouns or associate certain roles with specific demographics.

**Contamination risk:**

- Medium — Some overlap possible
- Explanation: BOLD's 23,679 prompts are derived from English Wikipedia, which is present in most training corpora. However, BOLD is not a benchmark that models are typically fine-tuned against explicitly, making contamination less severe than IFEval or TruthfulQA. The bias metrics (gender polarity, toxicity, sentiment) are distributional rather than accuracy-based, making direct memorisation less of a concern.

**Saturation risk:**

- Low — Benchmark is challenging
- Explanation: BOLD does not have a single accuracy score to saturate. It measures distributional properties of generated text across demographic groups, and models continue to show meaningful differences in bias across the profession and gender domains. There is no "perfect score" ceiling to hit.

**Format:**

- Free-form text (open-ended generation with automated bias metrics: sentiment, toxicity, gender polarity)

**Verdict:**

- Adapt it
- Explanation: BOLD's profession and gender subsets are directly usable for our scenario. We would use JD-style prompts referencing different professional roles and measure gender polarity and toxicity in the output. One limitation: BOLD uses a binary gender model and is US-centric in its racial categories. We would supplement it with manual review for non-binary and cross-cultural inclusivity, which BOLD alone cannot cover.