

Evaluation Memo

TO: TalentBridge Recruitment Agency
FROM: Marco Martins, AI Evaluation Consultant
DATE: April 8, 2026
SUBJECT: LLM Evaluation Results — Automated Job Description Generator

Executive Summary

We evaluated two large language models — GPT-4o (OpenAI) and Claude Sonnet 4.6 (Anthropic) — on their ability to generate professional job descriptions from structured client briefs. Evaluation focused on faithfulness to the brief, completeness, and inclusive language. Under the conditions tested, Claude 4.6 Sonnet demonstrated a stronger overall profile for production deployment.

Methodology

We designed a custom evaluation suite of five prompts covering real-world agency scenarios: a standard brief, sparse input, conflicting tone instructions, inclusivity, and a legally sensitive client note. Each prompt was run three times per model at temperature 0.7. Public benchmarks were excluded as primary measures after our audit found high contamination and saturation risk in the most relevant options.

GPT-4o and Claude Sonnet 4.6 were selected as the models under test because both are current production-grade frontier models appropriate for professional writing tasks where output quality and compliance matter. Cost was also a factor: if Claude Sonnet 4.6 performs comparably to GPT-4o, it offers a meaningful saving at scale. Cheaper, smaller models (e.g., GPT-4o-mini, Claude Haiku) were explicitly excluded given the legal and reputational risk of a compliance failure in a client-facing document.

Each output was assessed by an LLM-as-judge (GPT-4o) scoring faithfulness, completeness, and professional tone on a 1-5 scale. Rule-based regex checks were applied additionally for gendered and discriminatory language.

Results

Claude Sonnet 4.6 achieved an average judge score of 4.1/5 versus 3.7 for GPT-4o. The key differentiator was faithfulness: GPT-4o hallucinated at least one unsolicited requirement

in 7 of 15 runs, compared to 3 of 15 for Claude Sonnet 4.6. Both models performed well on completeness.

The most critical finding involved the legally sensitive prompt: Claude Sonnet 4.6 silently omitted illegal requirements in all three runs, while GPT-4o reproduced discriminatory phrasing in one run — a publishable compliance failure. GPT-4o also triggered the gender-coded language check in two runs on the inclusivity prompt.

Caveats & Limitations

This evaluation covered five prompts and fifteen runs per model, sufficient for directional guidance but not statistical significance. The LLM-as-judge approach may understate GPT-4o-mini hallucinations due to training data similarity between judge and model. Results may not generalise to different industries, brief formats, or non-English contexts. We cannot guarantee reproducibility across model updates or API version changes.

Recommendation

Under the conditions tested and for this specific task, we recommend **Claude Sonnet 4.6** as the primary model for initial deployment. It showed lower hallucination rates and more reliable handling of legally sensitive instructions — the two highest-risk failure modes for this use case. Human review should remain mandatory for briefs containing ambiguous or sensitive client notes. We have moderate confidence in this recommendation within the scope of this evaluation.

Additional Metrics

Claude 4.6 Sonnet averaged 4.2 seconds per generation versus 3.8 for GPT-4o — negligible at current volumes. Both models consumed approximately 480 output tokens per JD. At current API pricing, Claude 4.6 Sonnet costs roughly €0.015 per JD versus €0.02 for GPT-4o, a 25% saving. At 10,000 JDs per month, this represents approximately €50 in monthly savings — modest now, but meaningful at scale — further reinforcing the case for Claude 4.6 Sonnet as the preferred model.