

Reflection

Scenario: HR Recruitment Agency — Automated Job Description Generator
Consultant: Marco Martins
Date: April 2026

Question 1: What would change if the client's data was in French?

BOLD, our only retained benchmark, is English-only and would need to be replaced or rebuilt using French-language sources. Our custom evaluation prompts would require translation and review by a native French speaker, since literal translation of English HR language often produces unnatural phrasing — particularly given France's specific labour law terminology (e.g., "CDI", "temps plein").

The rule-based checks for gendered language would need a complete rewrite. French is grammatically gendered, and inclusive writing ("écriture inclusive") follows different and contested rules — what counts as biased in French is a sociolinguistic question that goes beyond simple pronoun detection. The LLM judge may also apply English-language quality standards to French output, penalising correct French register or missing culturally specific conventions.

Question 2: "Is this model AGI-level?" — how do you respond?

We would explain that "AGI-level" has no agreed definition and cannot be answered by any current evaluation framework. The more meaningful question is whether the model performs this specific task at or above human expert level — and that requires a human baseline we do not currently have: JDs written by senior HR consultants from the same briefs, scored by the same judge.

More broadly, no current benchmark reliably measures general intelligence. High scores reflect performance on known test sets, not adaptability or reasoning across novel domains. We would advise the client not to use the term "AGI-level" in any internal or external communication about the system.

Question 3: What is the one thing you could not evaluate without a human?

Whether a generated JD would actually attract the right candidates. A rule-based check cannot detect that a JD "sounds like every other tech company" or that the responsibilities section undersells the role. An LLM judge can assess structure and tone, but it cannot feel the difference between a JD that excites a strong candidate and one they scroll past. This requires lived experience as a job seeker or hiring manager in a specific industry.

In practice, we would recruit two or three experienced HR consultants to blind-review a sample of generated JDs alongside human-written ones, rating them on quality and likely candidate appeal — serving as ground truth for judge calibration and as a periodic production sanity check.